

Bayesian Optimisation of a Bio-Based Biodegradable Polyester Thermoset

Thomas Højlund-Dodd, James Elliott

Department of Materials Science and Metallurgy, University of Cambridge, Charles Babbage Road, Cambridge, CB3 0FS, UK

Abstract

Glass fibre-reinforced polymer composites are widely used in the wind turbine manufacturing industry, and are commonly manufactured using general-purpose unsaturated polyester resin thermosets. However, these resins are obtained from oil-derived aromatic monomers making them both unsustainably sourced and difficult to recycle. Bio-based and biodegradable alternatives exist but have not yet seen widespread commercial success due to their relatively high cost and poor mechanical properties. This paper seeks to address the question of whether bio-based branching thermoset polyesters can be optimised using Bayesian optimisation such that properties correlated with mechanical performance, including extent of reaction, might be improved upon. This is achieved by selecting a bio-based and biodegradable polymer (poly(glycerol citrate itaconate)), identifying a desirable objective variable (mass loss under esterification as proxy for polymerised content) and some tunable predictor variables (monomeric stoichiometry), and then applying sequential Bayesian techniques to uncover well-performing sample in an experimentally efficient manner. A Bayesian technique employing the ‘probability of improvement’ acquisition function discovered the experimental global optimum stoichiometry producing a mass loss of 20.77% in curing polyester samples. This represents 72% of the theoretical maximum, indicating a high extent of polymerisation. Meanwhile, an alternate sequential Bayesian technique deploying the ‘expected improvement’ acquisition function displayed high sample efficiency, outperforming the strongest benchmark set by a one-shot grid sampler in 63% of the samples. Both these sequential Bayesian techniques were simulated 1,000 times by querying an emulation of the practical experiment, respectively exhibiting a probability of discovering the hypothetical ‘region of global optimality’ of 1.1% and 5.3%. We conclude that Bayesian optimisation techniques represent a practically deployable and sample-efficient means by which desirable properties of bio-based thermosets may be improved.

1 Introduction

The early 21st century has seen continued growth in both the global human population (from 6.17 to 8.09 billion people) and total energy supply required to sustain them (from 416 to 634 EJ) between 2000 and 2023^{1–3}. In spite of continued strong reliance upon coal, natural gas, and oil for the provision of this demand, concerns regarding global climate and the finite nature of fossil reserves have galvanised investment in alternative renewable power systems^{3,4}. Of

these renewables, wind power saw the greatest growth in capacity to generate electricity (+2.33 PW h) ahead of both hydroelectric (+1.69 PW h) and solar photovoltaics (+1.60 PW h)⁵. Wind turbine blades are commonly manufactured from a composite comprised of thermosetting polymer matrix and reinforcing material, these strong thermoset matrices are chemically crosslinked which makes them difficult to reprocess or remould⁶. Given the relatively low rate of recycling and a common rate of replacement (ca. every fifteen years in China⁷) the in-

dustry generates a large volume of practically unrecyclable thermoset waste which contravenes the UN’s 12th SDG to “ensure sustainable consumption and production”⁸. In addressing the industry’s sustainability shortcoming identified above, the literature is divided into two groups: firstly, papers which focus upon improving the recyclability of contemporary materials, whether by thermal, chemical, mechanical or high-voltage fragmentation methods^{9–11}; and secondly, papers aimed at the development of properties in future replacement materials that would facilitate recycling¹², such as introduction of chemically cleavable bonds into the thermoset polymers used¹³. The motivation for present work was to ascertain whether Bayesian optimisation could be used to sample efficiently improve the degree of polymerisation in a novel biodegradable bio-based thermoset.

The composite matrix used for wind turbine blades is comprised of a moulding thermoset, with the most prevalently manufactured pre-thermoset form being general-purpose unsaturated polyester resin (GP-UPR)¹⁴. These GP-UPRs are based on a mixture of a reactive diluent (e.g. styrene) and a linear aromatic unsaturated polyester (e.g. poly(propylene glycol maleate phthalate)), which are cured chemically or thermally to generate the final thermosetting matrix¹⁵. Although polyesters are known to exhibit limited biodegradability (due to their labile ester bonds being prone to both biotic and abiotic hydrolysis¹⁶) it is the case that polyesters with aromatic moieties are more resistant than aliphatic counterparts to such degradation¹⁷. Consequently, thermosets derived from GP-UPRs synthesised from the aromatic monomer phthalic acid¹⁵ can be reasonably expected to exhibit a high resistance to biodegradation; which was part of the reason for their original choice as a material. An archetypical example of an aliphatic polyester would be poly(glycerol sebacate), which uses the monomers glycerol (triol) and sebacic acid (dicarboxylic acid) to form a

branching polyester thermoset on thermal curing via polyesterification¹⁸. Given their natural biodegradability, in the present work we chose to focus upon aliphatic polyesters.

Another concern associated with GP-UPRs with respect to the UN’s 12th SDG is that of a material’s provenance. If the underlying chemistry of a polymer is petrochemical in origin, then any carbon liberated from the material at end-of-life (whether by combustion for energy or biodegradation in landfill) will release fossil carbon, contributing via an anthropogenic route towards global climate change. Since GP-UPRs are synthesised from the oil-derived feedstock monomers propylene glycol, maleic acid, and phthalic acid, the resultant thermosets contain endogenic (fossil-derived) carbon¹⁹. In the present work we aim to avoid these feedstocks and focussed instead upon a bio-based alternative, with inspiration drawn from a chemical system presented in the literature by Alberts, where a combination of glycerol and citric acid were cured to form a bio-based branched aliphatic polyester thermoset²⁰. We build upon this research by adopting the chemical system and adapting it to include a third monomer, itaconic acid, which provides a second carboxylic acid with a lower degree of functionality (dicarboxylic acid); this second lower-functionality molecule was expected to encourage the emergence of alternative thermoset networks. Notably, Alberts found that physical properties of the end thermoset poly(glycerol citrate) could be tuned by modifying the monomeric stoichiometries of the prepolymer resin synthesised²⁰. Given this, we identified mass loss as a variable of interest during the synthesis of poly(glycerol citrate itaconate) as it is a proxy for polymerised content within a sample, and sought to optimise this by tuning the stoichiometries of the prepolymer resins mixed. Polymerised content could be an indicator of the crosslink density of these samples; a characteristic, which, if maximised, yields improvements in mechanical properties such as ultimate

tensile strength, Young’s modulus, and glass transition temperature²¹. Note that the link between sample mass loss and polymerised content is predicated on various assumptions: negligible volatilisation of monomers during curing, a lack of side-reactions between monomers, and a negligible mass of water retained within the final cured matrices of samples.

Poly(glycerol citrate itaconate) is formed from its monomers by polyesterification, a process which is dependent upon variables such as temperature, stoichiometric ratios, and additive presence²⁰. Additionally, reaction times for rigid solid thermoset materials are on the order of days when uncatalyzed²² (catalysis commonly involves strong fuming acids which can corrode equipment²³). Given the high dimensionality and temporally expensive function evaluation exhibited by this family of polyesters, a sample efficient global optimisation technique, namely Bayesian optimisation, was selected to tackle this problem. Alternative global optimisation techniques were considered but ruled out; for example, the hill-climbing class gradient descent algorithm is reliant on derivatives of the objective function which would have been uncertain given the noisy function evaluations associated with physical systems as opposed to simulated ones²⁴. The population class evolutionary and particle swarm optimisation algorithms were suitable but avoided in favour of the more sample efficient surrogate class of algorithms. Surrogate class algorithms are variously formulated, but the three major forms are distinguished by the surrogate models they leverage: polynomial response surface model, kriging (also known as gaussian process regression) and radial basis functions²⁵. Kriging is preferable to response surface modelling since it makes no assumption of polynomial order, avoiding a priori bias in surrogate models generated. Furthermore, radial basis function-based surrogate models provide no estimates of uncertainty around the fitted function; functions and associated

uncertainties derived using kriging provides the means by which Bayesian optimisation effectively balances exploitation and exploration of the parameter space. Hence, for the development of a part-computational and part-laboratory based workflow for the optimisation of mass loss in poly(glycerol citrate itaconate) we decided to use Bayesian optimisation predicated on kriging.

Originally coined in the 1970s by Mockus et al.²⁶, Bayesian optimisation was later developed further in the 1990s by Jones et al. into a technique called the efficient global optimisation algorithm²⁷- this name and technique was thereafter adopted and popularised by the DoE community. More recently, the original name has seen a revival as the technique and several adaptations have been re-discovered and implemented in a range of contemporary communities including robotics²⁸, bioengineering²⁹, and machine learning³⁰. In materials science, Bayesian optimisation has been used to predict crystalline structures³¹, carry out sample-efficient optimisation of density functional theory simulations³², and multi-objective optimisation of stability and ion-conductivity in materials³³.

This paper begins by detailing the polymeric system considered, experimental means by which samples were manufactured and characterised, as well as the means by which a suitable parameter space was devised. Hereafter, we detail the algorithmic mechanism behind the Bayesian techniques used before further exploring the initial one-shot quasirandom sampling strategy, sequential Bayesian acquisition functions, and one-shot benchmark sampling routines applied during the course of this work. The method concludes with a brief overview of how the physical system was computationally emulated such that Bayesian techniques could be simulated repeatedly to gain statistical insights. The results section is split between elucidation of results associated with the practical optimisation of the physical system, and those which derived from repeated


simulation of optimisation routines carried out upon the emulated system. The paper concludes that the Bayesian sequential routines succeeded in improving the mass loss metric from 48% to 72% of the theoretical limit, whilst also statistically proving that the ‘probability of improvement’ and ‘expected improvement’ acquisition functions respectively exhibited a 1.1% and 5.3% probability of discovering values approaching the emulated global optimum; this starkly contrasted with the benchmark one-shot techniques, for which, it was virtually impossible to discover such samples.

2 Method

An integrated experimental-computational methodology was developed to optimise the properties of the poly(glycerol citrate itaconate) thermoset, supported by custom software enabling human-in-the-loop optimisation in a laboratory setting. Programs were written in the Python programming language and developed into modules, whereafter these were united into a single package, “BO4BT”, available from its Github repository (<https://github.com/ThomasDodd97/BO4BT>) alongside its dependencies. A .yaml configuration file is also provided which allows an identical Conda environment to be setup, allowing a user to run this code³⁴.

2.1 Chemical Background

The condensation reaction occurring between carboxyl and hydroxyl groups is an esterification reaction:



As opposed to the simple esterification shown in Equation 1, the desired polyester thermoset, poly(glycerol citrate itaconate), is the product of polyesterification between multi-functional monomers.

Polyesterification is a form of polycondensation reaction which takes place as follows: an initial ester is formed from two difunctional monomers, where one must at least be a diol and the other a diacid (alternatively, each may possess a carboxyl and a hydroxyl group), producing a molecule of water and an ester with two spare functional groups- these groups can subsequently react with more difunctional monomers causing polymerisation into a linear polyester (See Figure 1). When monomers with a functionality greater than two are introduced into a polyesterification reaction, then side chains can form along the primary backbone of the growing polyester, which causes the development of a branching polymeric structure (See Figure 2). When a side chain reaches another large growing polyester molecule and reacts with it a crosslink is said to have developed, and when widespread crosslinking occurs generating a 3D network of polyester molecules a thermoset polyester has been created. Figure 2 shows how the monomers employed in the synthesis of poly(glycerol citrate itaconate) give rise to this branching thermoset structure.

Importantly, during the polyesterification process water is generated as the main by-product, and hence it is assumed that virtually all mass loss between prepolymer and final cured thermoset would be associated with water evaporation. Given that each molecule of water produced corresponds to an ester bond formed, mass loss can be said to be a good proxy for extent of polymerisation within a sample²². Thus, this research identifies mass loss as a metric of interest and sought to maximise this during optimisation of poly(glycerol citrate itaconate) synthesis: increasing the extent of polymerisation and thereby improving the degree of crosslinking and associated strength of these bio-based thermosets.




Figure 1: Archetypical polyesterification between a diol (propylene glycol) and a diacid (succinic acid) forming a linear polyester. Functional groups participating in the reaction are highlighted: hydroxyl groups blue, and carboxyl groups orange. In the products, ester groups are blue and carbonyl groups and water are also orange in accordance with provenance.

2.2 Experimental Method

Synthesis of the poly(glycerol citrate itaconate) thermoset was achieved via an uncatalysed polyesterification reaction (See Figure 2). Acid catalysis was avoided to reduce corrosion damage to equipment, this ultimately reduced the rate at which this reaction could be carried out but did not prevent polycondensation³². The monomers glycerol (99+ %), citric acid (99.6 %), and itaconic acid (99+ %) used were supplied by Fisher Scientific International and manufactured by Acros Organics.

Thermoset synthesis was two-step: firstly, a homogenous prepolymer resin was prepared by stirring solid monomers (acids) together with liquid monomer (glycerol) atop a hotplate and raising the temperature to a maximum of 140 °C until the solids were completely dissolved. Note, 140 °C was adopted to accelerate dissolution without encouraging decarboxylation of the citric acid to itaconic acid²⁰. Secondly, prepolymer was cured into the desired thermoset by transferral to a silicone mould before being exposed to elevated temperatures (90 °C) for an extended period (up to a week) within a drying oven.

Preliminary experimentation showed that temperatures exceeding 100 °C, caused boiling of byproduct water within samples. Gaseous water from boiling was entrapped within the curing polyester creating bubbles within the final structure, resulting in unpredictable densities of the final material.

Given the unpredictable nature of samples generated in this fashion, drying oven temperatures were held below 90 °C to avoid this uncontrollable effect upon the experiments.

As defined by Equation 2, determination of mass loss associated with the synthesis of a thermoset sample (Δm) was obtained by measuring the difference between both the initial mass of prepolymer resin (m_p) and the final mass of cured thermoset (m_t).

$$\Delta m = m_p - m_t \quad (2)$$

Equation 3 defines sample mass loss (Δm) normalised into a percentage format ($\Delta m\%$), this allowed inter-comparison since individual samples had slightly varying initial prepolymer resin masses (m_p).

$$\Delta m\% = \frac{\Delta m \times 100}{m_p} \quad (3)$$

Taking into account an initial prepolymer mass of 3 g, a mass loss range spanning 0.2-0.6 g, and a mass balance repeatability of 0.01 g we could estimate the uncertainty with respect to $\Delta m\%$ using a Monte Carlo sampling exercise and some assumed distributions- this equated to 0.44 % points with $1\sigma = 0.01$ % points (see supplementary information in Section 8.1). Residence time within the drying oven for a sample was controlled by periodic measurement of sample mass: mass was plotted against time, an exponential decay curve fitted, and the sample removed from the oven once a predetermined rate of mass loss was reached





Figure 2: Polyesterification between glycerol, citric acid, and itaconic acid forming poly(glycerol citrate itaconate). Functional groups participating in the reaction are highlighted: hydroxyl groups blue, and carboxyl groups orange. In the products, ester groups are blue and carbonyl groups and water are orange in accordance with provenance.

(this equated to around $-0.028 \text{ \% h}^{-1}$ of the initial mass).

2.3 Optimisation Framing

Having selected percentage mass loss within a thermoset sample ($\Delta m_{\%}$) as a suitable target variable for optimisation, predictor variables were defined and constrained creating a suitable parameter space describing a full range of prepolymer monomeric compositions.

The polyesterification of monomers to polymer seen in Figure 2 is presented in Equation 4 as a parameterised stoichiometric equation linking monomeric reagents glycerol (G), itaconic acid (It), and citric acid (Ci) with the product polymer poly(glycerol citrate itaconate) (P(GCiIt)) and byproduct water. Equation 4 introduces two stoichiometric variables representing the moles of glycerol (n_G), citric acid (n_{Ci}) and itaconic acid (n_{It}), and a variable representing the unknown number of moles of water produced (v).



The polyesterification reaction can be considered to consist of two possible polymerisation reactions between the triol glycerol and either the triacid citric acid or the diacid itaconic acid. By setting the glycerol stoichio-

metric variable to unity ($n_G = 1$), Figure 3 shows how the two separate reactions can be considered as if they were stoichiometric ranges, where there exist stoichiometric bounds between which polymerisation is possible.

These stoichiometric bounds are used to directly constrain the stoichiometric variables; Equation 5 describes the lower (n_{Ci}^l) and upper (n_{Ci}^u) bounds placed on the citric acid stoichiometric variable, whilst Equation 6 shows those for lower (n_{It}^l) and upper (n_{It}^u) bounds on proportion of itaconic acid.

$$n_{Ci}^l \leq n_{Ci} \leq n_{Ci}^u = 0.5 \leq n_{Ci} \leq 2.0 \quad (5)$$

$$n_{It}^l \leq n_{It} \leq n_{It}^u = 1.0 \leq n_{It} \leq 2.0 \quad (6)$$

In addition to the basic polymerisable bounds imposed on the two stoichiometric variables, a further constraint must be imposed for a successful optimisation: it must be possible for permutations of prepolymer monomeric mixtures to allow one acid to entirely dominate the other whilst ensuring that a prepolymer will still polymerise- this constraint ensures that any polymerisable stoichiometric mixture of monomers is possible. To achieve this, a three-dimensional bounded domain (a.k.a. parameter space, A) was devised consisting of three fundamental predictor variables: firstly, a citric acid strength predictor variable (s_1) influencing the citric acid stoichiometric variable. Secondly, an itaconic acid strength predictor




Figure 3: Simplified chemical representations displaying the polymerisable range of the two different triol-acid combinations encountered during poly(glycerol citrate itaconate) synthesis. These ranges represent the stoichiometric intervals where di-functional and trifunctional monomers allow network formation; outside these bounds gelation may not occur. In the legend: “G” glycerol, “It” itaconic acid, “Ci” citric acid.


variable (s_2) influencing the itaconic acid stoichiometric variable. Thirdly, a balancing predictor variable (b_1) capable of trading off the two acid systems, influencing both the citric and itaconic acid stoichiometric variables. $A \subseteq \mathbb{R}^d$ describes the parameter space A as a multi-dimensional real of d dimensions, where $d = 3$ given our three predictor variables. All three predictor variables took constrained ranges of between 0 and 1 ($s_1 = [0, 1]$, $s_2 = [0, 1]$, $b_1 = [0, 1]$), and were linked mathematically to the stoichiometric variables (via equations 7 and 8).

$$n_{Ci} = (n_{Ci}^l + (n_{Ci}^u - n_{Ci}^l) \times s_1) \times b_1 \quad (7)$$

$$n_{It} = (n_{It}^l + (n_{It}^u - n_{It}^l) \times s_2) \times (1 - b_1) \quad (8)$$

In summary, any random set of predictor variable values drawn from their respective constrained ranges of 0 to 1 (e.g. $s_1 = 0.75$, $s_2 = 0.25$, $b_1 = 0.5$) could be used to calculate a unique permutation of stoichiometric variables (e.g. $N_{Ci} = 0.81$, $N_{It} =$

0.63) which, on mixing of a prepolymer in accordance with Equation 9 and curing, would give rise to a unique polymer and associated mass loss (e.g. $\Delta m\% = 30$).



The sequence of calculations linking abstract predictor variables (s_1, s_2, b_1) to the practical stoichiometric variables (n_{Ci}, n_{It}) and onwards to the target variable of mass loss ($\Delta m\%$), allows the expression of a black-box mapping between predictors and targets. This black-box mapping is described by $f : A \rightarrow \mathbb{R}$, where an unknown function (f) maps any point in our parameter space (A) to a single real number value ($\mathbb{R}$). Equation 10 more explicitly expresses the black-box function in terms of the our target variable and predictor variables at a given sample number (i). Note that sample number is an integer counter used to track the number of function evaluations carried out

during a sequential optimisation.

$$\Delta m_{\%} = f(\mathbf{x}) = f([s_1^i, s_2^i, b_1^i]) \quad (10)$$

Our optimisation problem is thus expressed in Equation 11 as the maximisation of the black-box function, where the global optimum set of predictor variables ($\mathbf{x}^*$) is sought.

$$\mathbf{x}^* = \arg \max_{\mathbf{x} \in A} f(\mathbf{x}) \quad (11)$$

2.4 Optimisation Algorithm

The sequential optimisation algorithm employed during this work is summarised by the flow diagram in Figure 4. Successful conclusion of the optimisation is controlled by an if-else clause comparing an initially specified permissible number of function evaluations (a.k.a. sample limit, p_{tot}) with the current sample number (i): when $i < p_{\text{tot}}$ the algorithm continues. A second if-else statement controls whether a dataset exists from which a Gaussian process prior might be built for Bayesian sampling. When available, the algorithm uses the Gaussian process to draw a pre-defined number of samples (p_{iter}) from the parameter space for evaluation. If the prior-forming dataset is lacking, an initial quasirandom (Sobol) sampling of the parameter space with a pre-defined number of deterministic samples (p_{init}) is initiated. Samples drawn from the parameter space are processed into stoichiometric variables ($n_{\text{Ci}}, n_{\text{It}}$), with which an experimenter can manufacture sample(s) before evaluating the target variables $\Delta m_{\%}$ value(s) and returning these to the algorithm. The algorithm seen in Figure 4 was implemented in the Python programming language, and is incorporated into the ‘BO4BT’ package described earlier³⁴.




Figure 4: Flow diagram depicting the sequential optimisation algorithm developed.

On calculation of stoichiometric variables during an iteration, the algorithm merges all new sets of predictors variable values into a growing predictor-matrix ($\mathbf{X}$), shown in Equation 12.

$$\begin{aligned} \mathbf{X} &= [\mathbf{x}^{i=1}, \mathbf{x}^{i=2}, \dots] \\ &= [[s_1^{i=1}, s_2^{i=1}, b_1^{i=1}], [s_1^{i=2}, \dots], \dots] \end{aligned} \quad (12)$$

On conclusion of an iteration, the algorithm merges all data associated with samples evaluated into an optimisation-governing dataset. Equation 13 presents this dataset ($\mathbb{D}$) as a matrix which consists of the predictor-matrix ($\mathbf{X}$) and an array of target

variable values ($\Delta \mathbf{m}_{\%}$).

$$\begin{aligned}\mathbb{D} &= [\mathbf{X}, f(\mathbf{X})] \\ &= [\mathbf{X}, \Delta \mathbf{m}_{\%}] \\ &= [[\mathbf{x}^{i=1}, \mathbf{x}^{i=2}, \dots], [\Delta m_{\%}^{i=1}, \dots]]\end{aligned}\quad (13)$$

2.5 Sampling Procedures

During a sequential optimisation two forms of sampling were employed: firstly, a deterministic space-filling sampling of a quasirandom nature was used to representatively sample the parameter space when initialising the optimisation. Subsequently, a Bayesian sampling technique was used to sample-efficiently guide the search for the global optimum x^* .

2.5.1 Initial Quasirandom Sampling

Optimisation routines were initiated by representatively sampling across the parameter space A ; achieved using the quasirandom Sobol sampling algorithm to draw a pre-set number (p_{init}) of samples from across A . Sobol sampling was employed as it provides the same or better space-filling properties across $\mathbb{R}^{1 \leq d \leq 40}$ domains as compared with pseudorandom or grid techniques³⁵. On evaluation of the initial representative set of samples ($[\mathbf{x}^{i=1}, \dots, \mathbf{x}^{i=p_{\text{init}}}]$) for the target variable values ($[\Delta m_{\%}^{i=1}, \dots, \Delta m_{\%}^{i=p_{\text{init}}}]$), an initial prior-forming dataset ($\mathbb{D}^{i=p_{\text{init}}}$) is produced. In practice, only one quasirandomly sampled initial prior-forming dataset ($\mathbb{D}^{i=8}$) was generated and used (visualised in Figure 5); this allowed a fair inter-comparison of the two Bayesian optimisation routines attempted. As an aside, a relatively smooth topography witnessed during preliminary experimentation in lower dimensions was the justification for the use of only 8 samples during initial quasirandom sampling.

Prior-Forming Dataset, $\mathbb{D}^i = 8$




Figure 5: Initial 8 samples quasirandomly from across the parameter space, A , comprised of s_1 , s_2 , and b_1 (as in equations 7 and 8), which provide the prior-initialising dataset for the Bayesian routines. Coloration of points denote distance from viewer with darker points appearing further away.

2.5.2 Sequential Bayesian Sampling

Sequential sampling by means of Bayesian optimisation was undertaken using a Gaussian process regression model, which was derived from the regression of an optimisation dataset $\mathbb{D}$. Equation 14 shows how the Gaussian process, $\mathcal{GP}$, is a distribution over functions which is completely described by both a mean function, m , and a covariance function, K .

$$f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), K(\mathbf{x}, \mathbf{x}')) \quad (14)$$

As an infinite-dimensional generalisation of multivariate normal distributions the Gaussian process inherits the properties of a conjugate prior³⁶, allowing the initial posterior (e.g. $\mathcal{GP}^{i=8}$) to be recycled as a prior alongside new data (e.g. $\mathbb{D} + [\mathbf{x}^{i=9}, \Delta m^{i=9}]$) to generate an updated posterior (e.g. $\mathcal{GP}^{i=9}$) in accordance with Bayes' theorem. The surrogate model evolves during sequential optimisation as new data are incorporated at each iteration, i .

During the fitting of the Gaussian process, a covariance function is used to encode beliefs regarding the families of functions likely to fit the data observed. In

this work we used a Matérn covariance function containing a hyperparameter controlling smoothness of functions drawn from the Gaussian process. This was set to $\nu = 2.5$, balancing the need for smoothness in sample paths against the drawbacks of higher values of ν (e.g. unphysical nature of infinitely differentiable functions or mathematical instabilities when undertaking likelihood calculations).

After generating a posterior, a simple one-step lookahead sampling policy was used to select a suitable next point $\mathbf{x}$ within the parameter space A to sample. In this work, we focused on policies making use of the ‘simple reward’ utility function (see Equation 15) which takes the current dataset $\mathbb{D}$, and from the set of sample coordinates evaluated so far X returns the predicted mean value $\Delta m\%$ of the point $\mathbf{x}$ exhibiting the highest predicted mean ($\max m(X)$) according to the current posterior $\mathcal{GP}^{37}$.

$$u(\mathbb{D}) = \max m(X) \quad (15)$$

In adopting this utility function, two different sampling policies were developed and investigated. The first was predicated on the ‘probability of improvement’ acquisition function (defined in Equation 16), where any given putative $\mathbf{x}$ is scored by the probability that it will produce an updated dataset $\mathbb{D}'$ which generates a higher evaluation from the utility function than some absolute threshold τ . The probability of a sample exceeding the current best is deduced using the current posterior $\mathcal{GP}$ ’s mean ($m(\mathbf{x})$) and standard deviation ($\sqrt{K(\mathbf{x}, \mathbf{x})}$) measured at that sample position. Here, the improvement threshold was defined as the current best observed value ($\tau = u(\mathbb{D})$), rather than introducing an additional offset parameter ε . This simplified the optimisation while avoiding unnecessary hyperparameter tuning.

$$\begin{aligned} \alpha_{PI}(\mathbf{x}) &= Pr(u(\mathbb{D}') > \tau) \\ &= \int_{\tau}^{\infty} \mathcal{N}(m(\mathbf{x}), \sqrt{K(\mathbf{x}, \mathbf{x})}) dm(\mathbf{x}) \quad (16) \\ &= 1 - \Phi((\tau - m(\mathbf{x}))/\sqrt{K(\mathbf{x}, \mathbf{x})}) \end{aligned}$$

The second policy made use of the ‘expected improvement’ acquisition function (defined in Equation 17), which incorporates an expectation of improvement value ($u(\mathbb{D}') - u(\mathbb{D})$). This counteracts any overtly exploitative behaviour by rewarding both the scale and probability of an improvement above the threshold τ^{38} .

$$\alpha_{EI}(\mathbf{x}) = \int_{\tau}^{\infty} c \mathcal{N}(m(\mathbf{x}), \sqrt{K(\mathbf{x}, \mathbf{x})}) dm(\mathbf{x}) \quad (17)$$

where $c = u(\mathbb{D}') - u(\mathbb{D})$

Equation 18 shows how the next point to be sampled, $\mathbf{x}^+$, is found by trivially optimising the cheap acquisition function (many putative samples $\mathbf{x}$ may be evaluated) to find the global optimum within the bounds of the parameter space A^{38} .

$$\mathbf{x}^+ = \arg \max_{\mathbf{x} \in A} \alpha(\mathbf{x}) \quad (18)$$

Gaussian processes and acquisition functions were implemented using two open-source packages: scikit-optimize (“skopt”, v0.9.0)³⁹ and Ax (“ax-platform”, v0.4.3)⁴⁰. Note that the Gaussian process regression models developed by these software packages assumed that measured percentage mass loss values comprised both an underlying composition-dependent response plus some independent and normally distributed noise. Both the Ax and scikit-optimize packages inferred the homoscedastic noise during the fitting of the Gaussian process, with σ values of 0.35% points and 0.46% points recorded respectively (this being for the final Gaussian processes obtained at culmination of the two optimisation routines undertaken).

2.6 Benchmarks

The sampling efficiency of the Bayesian sequential optimisation techniques was demonstrated by employing three simple benchmarks. These were derived using one-shot deterministic space-filling sampling techniques including pseudorandom, quasirandom (Sobol), and grid sampling to disperse the maximum allowable number of samples




Figure 6: *Deterministic space-filling one-shot sampling techniques. Coloration of points denote distance from viewer with darker points appearing further away.*

($p_{\text{tot}} = 27$) across the parameter space $\mathcal{A}$ (See Figure 6). Subsequently, for each of the techniques deployed, all samples, $\mathbf{X}$, were evaluated, $f(\mathbf{X})$, to retrieve associated target variable values, $\Delta m\%$, generating a complete dataset $\mathbb{D}$. The highest target variable value was identified within each technique’s dataset ($\Delta m_{\%}^{\circ}$) and equated to the benchmark value (See Table 1).

Benchmark Technique	Max (%)	Mean (%)	Standard Dev. (%)
Grid	16.24	11.93	1.66
Quasirandom	14.55	12.34	1.43
Pseudorandom	14.01	12.54	0.88

Table 1: *Statistical overview of mass loss, $\Delta m\%$, target variable evaluations obtained from each of the three 27-sample datasets generated by the one-shot deterministic sampling techniques used as benchmarks. ‘Max’ equates to the highest discovered mass loss, ($\Delta m_{\%}^{\circ}$), ‘Mean’ is the population mean, and ‘Standard Dev.’ relates to the standard deviation.*

2.7 Emulation of Physical System

Since repeated experimental optimisation was impractical, a computational emulator of the physical system was constructed to facilitate statistical comparison of the sampling strategies. The emulator was trained on a combined dataset of all experimentally evaluated samples from both the benchmark and sequential optimisation methods and was used as an approximate, data-driven, representation of the objec-

tive function within the region of parameter space explored experimentally. Subsequently, a Gaussian process $\mathcal{GP}$ was trained on this master dataset using the parameterisations described earlier- the resultant mean function $m(\mathbf{x})$ then being used as an inexpensive surrogate function of the latent objective function (i.e. the emulator). Note, that some data obtained outside the sample limit was included in this master dataset- specifically, the final samples generated by the Bayesian sequential technique using the ‘probability of improvement’ acquisition function. These were remarkably high, and represented the experimental global optimum discovered during the course of this work- as opposed to the true global optimum associated with the latent objective function, this term denotes the globally optimal result given the experimental data available. Care was taken to ensure the trained $\mathcal{GP}$ faithfully replicated this region of optimality within the parameter space as well as the general trends of the objective function.

3 Results

This work produced two sets of results: firstly, those derived from optimisation conducted via sampling directly from the latent objective function described by the physical system and secondly, those stemming from the optimisation of a surrogate function trained to emulate the latent objective

function and the associated physical system. The latter set of results were generated to provide statistical support for the former.

3.1 Physical System Optimisation

Two Bayesian sequential optimisation experiments were run, each with the same sample limit ($p_{\text{tot}} = 27$), number of samples drawn per Bayesian iteration ($p_{\text{iter}} = 3$), and a prior-forming dataset of 8 quasirandomly sampled targets (See Figure 5). The difference between these two methods was the acquisition function employed: the first using the ‘probability of improvement’ (Equation 16) whilst the second used the ‘expected improvement’ function (Equation 17). Figure 7 shows the optimisation progression of these two sequential techniques when compared with the one-shot deterministic methods; mass losses plotted on the y-axis are the best values uncovered given the samples evaluated so far, $\Delta m_{\%}$. A statistical overview is presented in Table 2 which concerns the 8 quasirandom and 16 sequential samples for each of the two Bayesian techniques. Given the sequential samples available ($p_{\text{tot}} - p_{\text{init}} = 19$) and the samples per iteration ($p_{\text{iter}} = 3$), an extra iteration of sequential sampling was run which consequently breached the sample limit- these extra samples are included in Figure 7, but are not considered in the preparation of Table 2 (as it would be unfair to compare two different sampling techniques which were allowed an unequal number of samples over which to optimise).

By comparing Tables 2 and 1, it is apparent that both sequential routines outperformed the one-shot techniques in terms of best sample discovered.

Figure 7 shows that the sequential Bayesian technique employing the ‘expected improvement’ (EI - solid green line) acquisition function discovered a point which makes a statistically significant improvement upon the strongest one-shot benchmark ex-

ceeding the strongest one-shot technique in almost half the allowable samples (63% p_{tot}). In contrast, the ‘probability of improvement’ (PI - solid blue line) acquisition function failed to make a statistically significant improvement upon the grid sampled benchmark within the sample limit, although the sample drawn on the sample limit (100% p_{tot}) did make significant improvements upon both the pseudo/quasirandom benchmarks; this was consistent with the ‘probability of improvement’ acquisition function exhibiting more exploratory behaviour than the ‘expected improvement’ acquisition function. Further evidence of the ‘probability of improvement’ acquisition function’s exploratory bias is seen in Table 2, where the samples generated by this sequential routine obtain a markedly lower mean than that of those drawn by the ‘expected improvement’ acquisition function; anecdotally, on reviewing the placement of samples in 3D space it was noted that these were generally placed in exploratory locations within the parameter space. Contrary to our initial expectations, the evidence presented above shows that the ‘probability of improvement’ acquisition function does not exhibit overly exploitative behaviour due to the lack of a marginal threshold parameter ϵ being manually tuned.

Acquisition Function	Max (%)	Mean (%)	Standard Dev. (%)
PI	16.67	12.50	1.63
EI	17.99	13.67	1.91

Table 2: Statistical overview of mass loss, $\Delta m_{\%}$, target variable evaluations obtained from the two 27-sample sequential datasets: these either utilising the EI (‘expected improvement’) or PI (‘probability of improvement’) acquisition functions during Bayesian optimisation. ‘Max’ equates to the highest discovered mass loss, ($\Delta m_{\%}^{\circ}$), ‘Mean’ is the population mean, and ‘Standard Dev.’ relates to a single standard deviation.

Notably, the Bayesian sequential technique employing the ‘probability of improvement’ acquisition function discovered samples attaining the highest mass losses of




Figure 7: A plot comparing the optimisation progression for the two Bayesian sequential techniques alongside the three one-shot benchmark techniques. Expected improvement (EI) - solid green line. Probability of improvement (PI) - solid blue line. One-shot techniques - dashed lines. Uncertainty in $\Delta m_{\%}$ values is equal to 0.44% points (1σ).

any sample characterised during this work ($\Delta m_{\%}^{\circ} = 20.77\%$), albeit outside the sample limit (See Figure 7).

3.2 Emulated System Optimisation

Since curing of a single prepolymer sample could take up to 2 weeks to complete, the temporal expense of sampling the objective function was too high to allow for repeat experiments involving the sequential techniques deployed. This precluded statistically proving the utility of Bayesian optimisation in this field within reasonable timeframes. However, by developing an inexpensive surrogate function capable of emulating the physical system, a much larger number of function evaluations could be carried out computationally, facilitating statistical analysis of sampling methods' efficacy. By despatching 1,000 simulations of each of the two random benchmark one-shot techniques and the two Bayesian sequential techniques, and returning the best target variable values $\Delta m_{\%}^{\circ}$ discovered by each of these, basic statistical conclusions could be drawn. Note that each simulation of the sequential techniques was initiated with a separate quasirandomly drawn prior of 8 samples and that

the grid sampler was deprecated since its outcomes were not affected by chance.

Figure 8 presents a swarmplot showing distribution of outcomes for the 4,000 simulations across the four sampling techniques studied. On comparison of the mean $\Delta m_{\%}^{\circ}$ values discovered by the various sampling techniques, sequential Bayesian techniques were shown to generally perform better (PI=14.6%, EI=15.4% for $\Delta m_{\%}^{\circ}$) than one-shot techniques (PR=14.0%, QR=14.1% for $\Delta m_{\%}^{\circ}$). Although standard deviation values associated with these mean values would preclude a straightforward conclusion that the difference between values are statistically significant, it can be seen that the simulations are distributed in a somewhat lognormal fashion which shows that only the Bayesian techniques are capable of discovering well-above-average-performing samples (causing the majority of skew seen in the means). The Bayesian sequential technique adopting the 'expected improvement' acquisition function in its sampling policy not only outperformed the 'probability of improvement' technique in terms of mean value attained, but also in number of simulations obtaining $\Delta m_{\%}^{\circ}$ values which resided within the small region of global optimality. This region was characterised as containing function evaluations ex-




Figure 8: Swarm plot of maximum function evaluations obtained, $\Delta m_{\%}^{\circ}$, for each of the 4,000 simulated optimisations carried out. All optimisations used 27 samples, each technique was simulated 1,000 times, and the Bayesian routines were initiated with 8 quasirandomly sampled points.

ceeding 18.5%, and this simple threshold was therewith used to calculate the number of simulations attaining such globally optimum values ($\Delta m_{\%}^{\circ} > 18.5$) for each of the sampling techniques (PI=11, EI=53 simulations where $\Delta m_{\%}^{\circ} > 18.5$). Notably, neither of the one-shot techniques was able to uncover the globally optimum region (PR=0, QR=0 simulations where $\Delta m_{\%}^{\circ} > 18.5$). Given the above, and a (particularly) tentative assumption that the final emulation was an accurate description of the underlying objective function, we could venture that the probability of having discovered the experimental global optimum using the ‘expected improvement’ and ‘probability of improvement’ techniques was 5.3% and 1.1% respectively. This would make the practical reality of having discovered these points remote but most definitely not impossible; and, importantly, the probabilities are a marked improvement above the 0.0% exhibited by the one-shot techniques.

4 Conclusion

Having engaged in the practical optimisation of the monomeric composition of a simple bio-based branching thermoset polyester (poly (glycerol citrate itaconate)) with respect to an objective of interest (percentage mass loss via curing) we have discovered an experimental global optimum of $\Delta m_{\%}^{\circ} = 20.77\%$ at the sample coordinates ($s_1 = 0.337, s_2 = 0.648, b_1 = 0.773$) corre-

sponding to a stoichiometric balance of 1 glycerol to 0.77 citric acid to 0.37 itaconic acid. This being discovered by a sequential Bayesian technique whose sampling policy adopted the well known ‘probability of improvement’ acquisition function. Furthermore, this sequential Bayesian technique and a second (which employed the alternative ‘expected improvement’ acquisition function) displayed better sample-efficiency than simple one-shot techniques in the field.

Utility of these Bayesian techniques in sample-efficient global optimisation was proven through 1,000 simulations of their behaviour when sampling a surrogate function emulating the objective function; the ‘expected improvement’ and ‘probability of improvement’ sequential techniques were found to exhibit a 5.3% and 1.1% probability of uncovering the experimental region of global optimality, whilst both the pseudorandom and quasirandom one-shot techniques entirely failed to sample these regions.

If the intention is for poly(glycerol citrate itaconate) to perform the role of a useful thermoset resin for moulding applications, future work would have to focus on the mitigation of the bubbling effect, improvement of mechanical strength, and reductions in the cure time. The last of these is of especial importance given the relatively higher amounts of process-energy required to prepolymerise and cure the presently described branched thermosets when compared with conven-

tional unsaturated polyester resins (UPR). [Project Reference EP/S022961/1].

Modelling of these two processes with error propagation as described in the supporting information (see Section 8.2) suggested that conventional UPR thermosets require 81% less process energy than the branched-polyester route, with a 5th–95th percentile uncertainty range of 57–91%. This difference was found to be dominated by the prolonged 90°C oven residence time required for the branched polyester. The estimates remained sensitive to oven power, batch loading, cure duration and the process boundaries applied to prepolymer manufacture.

We therefore suggest future work perhaps focus upon dual-objective optimisation problems which consider both mechanical strength maximisation alongside oven residence time minimisation, whilst adopting predictor variables such as prepolymer stoichiometry, temperature & pressure, and quantities of catalytic reagent added. Alternatively, different bio-based and readily-available monomers could be considered (e.g. succinic acid) which may encourage development of novel thermoset networks that might impact upon the objectives of interest. Note that addition of further monomers to the stoichiometric balancing mechanism described in this paper would, due to the curse of dimensionality, require many more samples to effectively explore and optimise these systems: this would encourage the replacement of human-in-the-loop experimentation with laboratory automation.

5 Conflicts of interest

There are no conflicts to declare.

6 Acknowledgements

This work was supported by the Engineering and Physical Sciences Research Council as a part of the UKRI centre for doctoral training in the ‘application of artificial intelligence to the study of environmental risks’

7 Bibliography

1. In Data), O. (W. *Population Growth* (eds Ritchie, H. *et al.*) Retrieved October 27th, 2025. 2025. <https://ourworldindata.org/population-growth>.
2. Nations), U. (*World Population Prospects 2024* Retrieved October 27th, 2025. 2025. <https://population.un.org/wpp/downloads?folder=Standard%20Projections&group=Most%20used>.
3. IEA. *Energy Statistics Data Browser; Total Energy Supply by Source* Retrieved 16th October 2025. International Energy Agency, 2025. <https://www.iea.org/data-and-statistics/data-tools/energy-statistics-data-browser?country=WORLD&fuel=Energy%20supply&indicator=TESbySource>.
4. On Climate Change, I. P. *Climate Change 2021: The Physical Science Basis. Contribution of Working Group I to the Sixth Assessment Report of the Intergovernmental Panel on Climate Change* (eds Masson-Delmotte, V. *et al.*) (Cambridge University Press, Cambridge, 2021).
5. IEA. *Energy Statistics Data Browser; Electricity generation by source* Retrieved 27th October 2025. International Energy Agency, 2025. <https://www.iea.org/data-and-statistics/data-tools/energy-statistics-data-browser?country=WORLD&fuel=Energy%20supply&indicator=ElecGenByFuel>.
6. Ma, S. *et al.* Degradable thermosets based on labile bonds or linkages- A review. *Progress in Polymer Science* **76**, 65–110 (2017).

7. Chen, J. *et al.* Recycling and reuse of composite materials for wind turbine blades- An overview. *Reinforced Plastics and Composites* **38**, 567–577 (2019).
8. UN. *The 2030 Agenda for Sustainable Development* technical (United Nations, 2015).
9. Utekar, S. *et al.* Comprehensive study of recycling of thermosetting polymer composites - Driving force, challenges and methods. *Composites Part B* **207** (2020).
10. Beauson, J. *et al.* Recycling of shredded composites from wind turbine blades in new thermoset polymer composites. *Composites- Part A* **90**, 390–399 (2016).
11. Chen, R. *et al.* Comparative pyrolysis characteristics of representative commercial thermosetting plastic waste in inert and oxygenous atmosphere. *Fuel* **246**, 212–221 (2019).
12. Liu, Y. *et al.* Closed-loop chemical recycling of thermosetting polymers and their applications- a review. *Green Chemistry* **24**, 5691–5708 (2022).
13. Tesoro, G. *et al.* Reversible Crosslinking in Epoxy Resins. I. Feasibility Studies. *Applied Polymer Science* **39**, 1425–1437 (1990).
14. Bobade, S. *et al.* Bio-based thermosetting resins for future generation- a review. *Polymer-Plastics Technology and Engineering* **55**, 1836–1896 (2016).
15. Ahamad, A. *et al.* Synthesis of unsaturated polyester resin - effect of anhydride composition. *Designed Monomers and Polymers* **4**, 261–268 (2001).
16. Sharma, S. *et al.* *A handbook of applied biopolymer technology* (RSC Publishing, Cambridge, 2011).
17. Müller, R. *et al.* Biodegradation of polyesters containing aromatic constituents. *Biotechnology* **86**, 87–95 (2001).
18. Zhang, H. *et al.* Recent advances in glycerol polymers- chemistry and biomedical applications. *Macromolecular Rapid Communications* **35**, 1906–1924 (2014).
19. Levi, P. *et al.* Mapping global flows of chemicals- from fossil fuel feedstocks to chemical products. *Environmental Science and Technology* **52**, 1725–1734 (2018).
20. Alberts, A. *et al.* Plantics GX- a biodegradable and cost-effective thermoset plastic that is 100 percent plant-based. *Faraday Discussions* **202**, 111–120 (2017).
21. Panic, V. *et al.* Simple one-pot synthesis of fully biobased unsaturated polyester resins based on itaconic acid. *Biomacromolecules* **18**, 3881–3891 (2017).
22. Trenkel-Amoroso, J. *Synthesis, degradation and practical applications of a glycerol-citric acid condensation polymer* phd (Oregon State University, 2008).
23. Shukla, K. A study on the synthesis of various polyesters from glycerol. *Journal of Polymer Research* **29** (2022).
24. Stork, J. *et al.* A new taxonomy of global optimization algorithms. *Natural Computing* **21**, 219–242 (2022).
25. Roeva, O. *Real-world Applications of Genetic Algorithms* ISBN: 978-953-51-0146-8 (InTech, 2012).
26. Mockus, J. *et al.* The application of bayesian methods for seeking the extremum. *Towards Global Optimisation* **2**, 117–128 (1978).

27. Jones, D. *et al.* Efficient global optimization of expensive black-box functions. *Journal of Global Optimization* **13**, 455–492 (1998).
28. Calandra, R. *et al.* Bayesian optimization for learning gaits under uncertainty. *Annals of Mathematics and Artificial Intelligence* **76**, 5–23 (2016).
29. Bader, J. *et al.* Improving extracellular vesicles production through a Bayesian optimization-based experimental design. *European Journal of Pharmaceutics and Biopharmaceutics* **182**, 103–114 (2023).
30. Cho, H. Basic enhancement strategies when using Bayesian optimization for hyperparameter tuning of deep neural networks. *IEEE Access* **8**, 52588–52608 (2020).
31. Yamashita, T. *et al.* Crystal structure prediction accelerated by Bayesian optimization. *Physical Review Materials* **2** (2018).
32. Balachandran, P. *et al.* Adaptive strategies for materials design using uncertainties. *Scientific Reports* **6** (2016).
33. Karasuyama, M. *et al.* Computational design of stable and highly ion-conductive materials using multi-objective Bayesian optimization- case studies on diffusion of oxygen and lithium. *Computational Materials Science* **184** (2020).
34. Højlund-Dodd, T. *BO4BT* version 1.0.0. June 6, 2024. <https://github.com/ThomasDodd97/BO4BT>.
35. Garud, S. *et al.* Design of computer experiments- A review. *Computers and Chemical Engineering* **106**, 71–95 (2017).
36. Deisenroth, M. *et al.* *Mathematics for Machine Learning* (Cambridge University Press, 2020).
37. Garnett, R. *Bayesian Optimization* (Cambridge University Press, 2023).
38. mhstrongman. *advml2020* version 0.0.1. Jan. 1, 2020. <https://github.com/mhstrongman/advml2020>.
39. Kernc. *scikit-optimize* version 0.9.0. June 6, 2021. <https://github.com/scikit-optimize/scikit-optimize>.
40. Facebook. *Ax* version 0.4.3. June 6, 2024. <https://github.com/facebook/Ax>.

8 Supporting Information

Here we include some peripheral elements of work that support the paper.

8.1 Notes on Measurement Error Propagation

To evaluate uncertainty associated with the percentage mass loss metric we undertook some error propagation analysis by modelling the process of obtaining values for this metric. This modelling was carried out in the Python programming language and is detailed in a Jupyter notebook held within the Github repository associated with this paper, it can be found under the ‘Sandbox/UncertaintyProagation’ directory and is entitled ‘UncertaintyModel.ipynb’.

By setting the initial mass of a given sample at 3 g and allowing the final mass to vary in accordance with a uniform distribution between a lower and upper bound (as deduced by from the two samples evaluated exhibiting the maximal and minimal mass losses across all trials) we could simulate the drawing of a random sample from across the parameter space each having its own unique mass loss. Additionally, to model the measurement error which manifests itself during the process of evaluating the initial mass and final mass values we could employ normal distributions with standard deviations

equal to the repeatability of the scales employed (0.01 g). By coupling these concepts together and undertaking a Monte Carlo sampling exercise we could obtain estimates of the average uncertainty associated with any given percentage mass loss value considered in the present work; this was equal to 0.44 % points ($1\sigma=0.01\%$ points).

8.2 Notes on Process-Energy Calculations

Although this paper has focussed upon branching polyester thermosets as an alternative to unsaturated polyester resin derived thermosets on account of the relative ease with which these materials may be prepared from bio-based feedstocks, we acknowledge the fact that this does not inherently make this family of polymer more sustainable than their conventional counterparts. To prove this, one would first have to choose a suitable quantitative definition of sustainability (e.g. embodied carbon), before engaging in a comprehensive life-cycle analysis for both the methods by which thermosets were manufactured.

Although by no means a full life-cycle analysis of either route, we have undertaken some modelling of the two system with a view to elucidating whether the prolonged thermal cure time used for branched poly(glycerol citrate itaconate) thermosets is likely to impose an energy penalty relative to the unsaturated

polyester resin route. This model is presented within a Jupyter notebook file entitled ‘calcs_restructured_uncertainty.ipynb’ (within the ‘Sandbox/EnergyCalculations’ directory) in the github repository associated with this paper.

Here, we establish two separate models, one for the branched system and the other for the unsaturated polyester system. These models both consider the energy required to form prepolymers from feedstocks and the subsequent energy required to cure these prepolymers. Each model is based on a set of parameters, each parameter being described by a triangular distribution, which allows propagation of uncertainty to be considered over the course of a large number of simulations in a Monte Carlo fashion.

On conclusion of 20,000 Monte Carlo draws, it was concluded that the more energy intensive route was that of branched polyester thermoset manufacture (requiring an average of 5.4x the kWh kg⁻¹ of thermoset). Table 3 provides an overview of this exercise, and shows how the propagated error bounds surrounding the median values. Estimates of kWh kg⁻¹ yielded by the branched model were particularly sensitive toward power required to maintain the oven temperature as well as the cure duration, whilst the unsaturated polyester model was most sensitive toward energy required to maintain reactor temperature, auxiliary power requirements, and heat capacity of various elements.

Metric	5 th Percentile	Median	95 th Percentile
Branched Polyester Route (kWh kg ⁻¹)	24.4	49.9	84.1
Conventional UPR Route (kWh kg ⁻¹)	5.4	9.2	14.6
UPR/Branched Ratio	0.087	0.186	0.431

Table 3: *Summary of process-energy modelling exercise.*